


CSCI 544: Applied Natural Language Processing

Advanced Techniques in Large-Scale Pretraining

Xuezhe Ma (Max)



USC University of
Southern California


Large-Scale Pretraining

- **Model Architecture**
 - Transformer
- **Data**
 - Publicly available data
 - Common Crawl
 - Wikipedia
 - Books
 - ...
- **Training Objective**
 - Autoregressive language modeling (next-token prediction)
- **Machines**
 - Sufficient amount of GPUs
 - A100 (80G memory per device)

Training Receipt

- **7B parameters**
 - 32 blocks of Transformer decoder
 - Model dimension size $d = 4096$
- **Training sequence length**
 - 4096 tokens
- **Training batch size**
 - 4 million tokens (1,000 sequences)

Challenges in Large-Scale Pretraining



16 billion elements (64G memory w. float32)
distributing the large tensor to multiple GPUs!

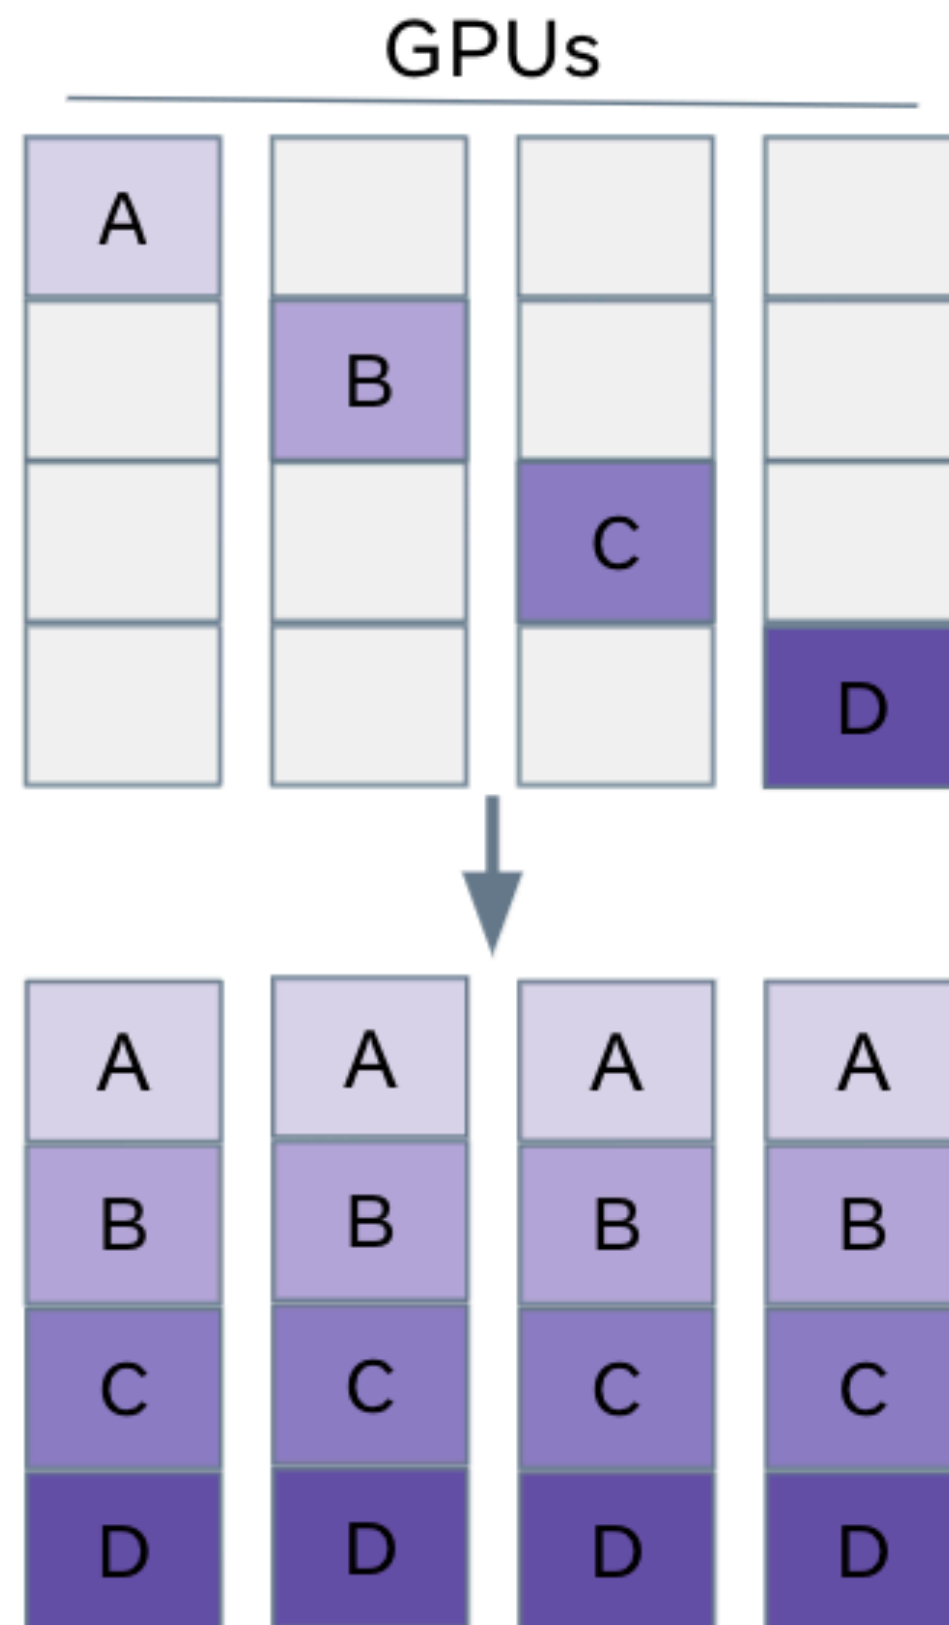
Communication between GPUs

- For distributed training, we need to synchronize **training status**
 - Model parameters
 - Gradients
 - Optimization states

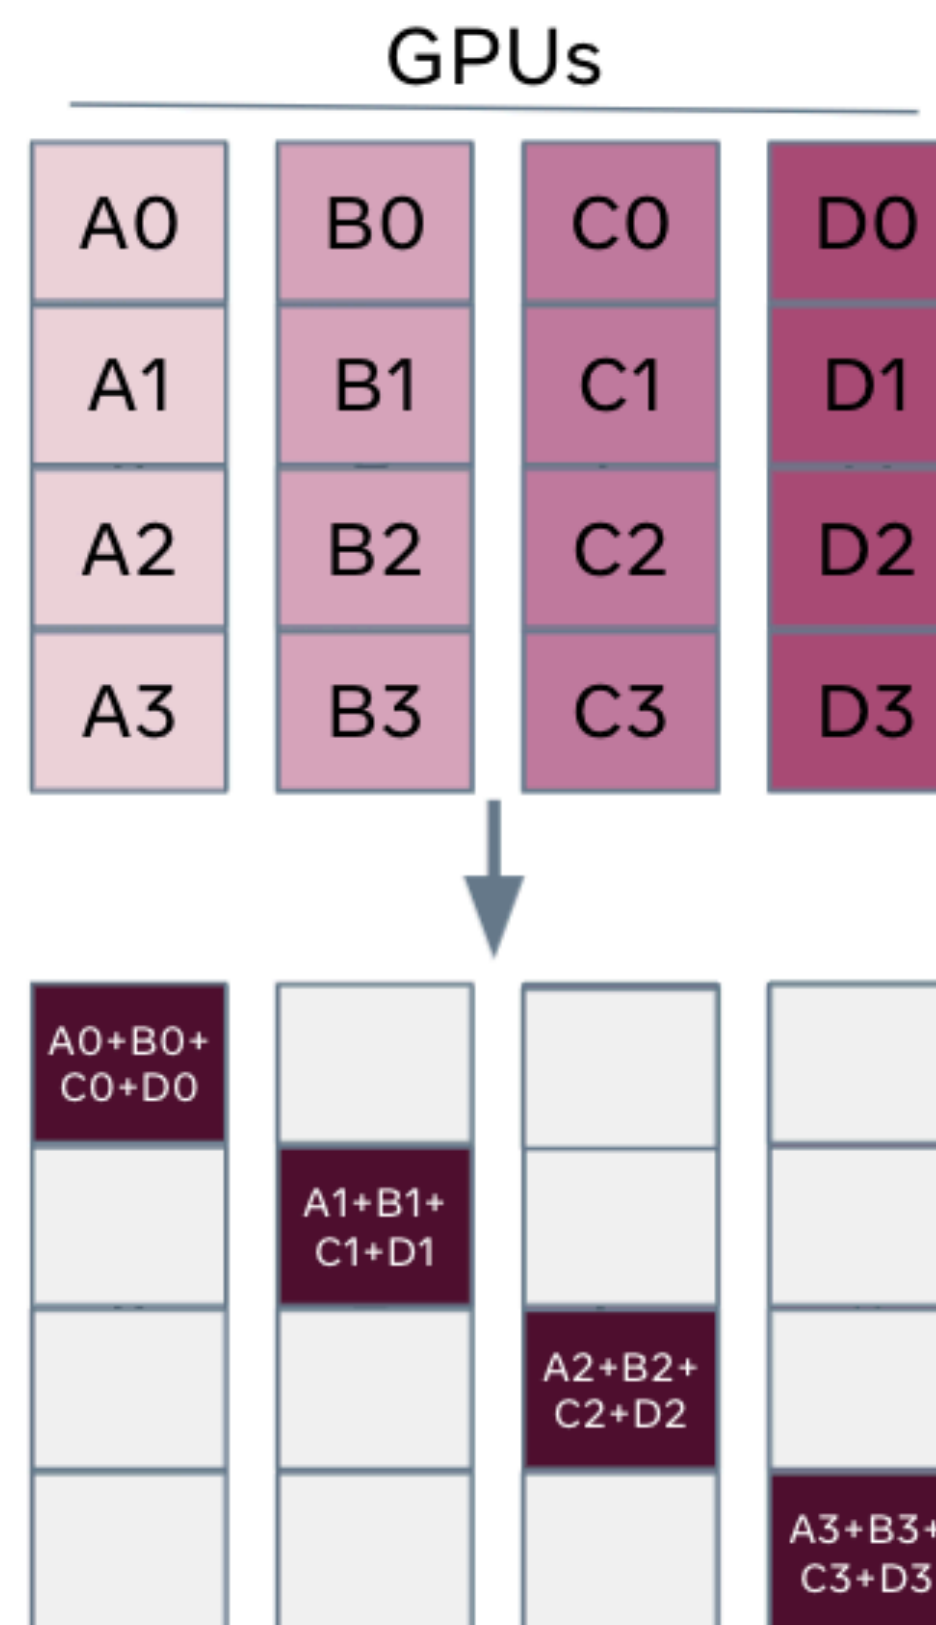
Communication between GPUs

- Some basic communication operations


All-gather



Reduce-scatter

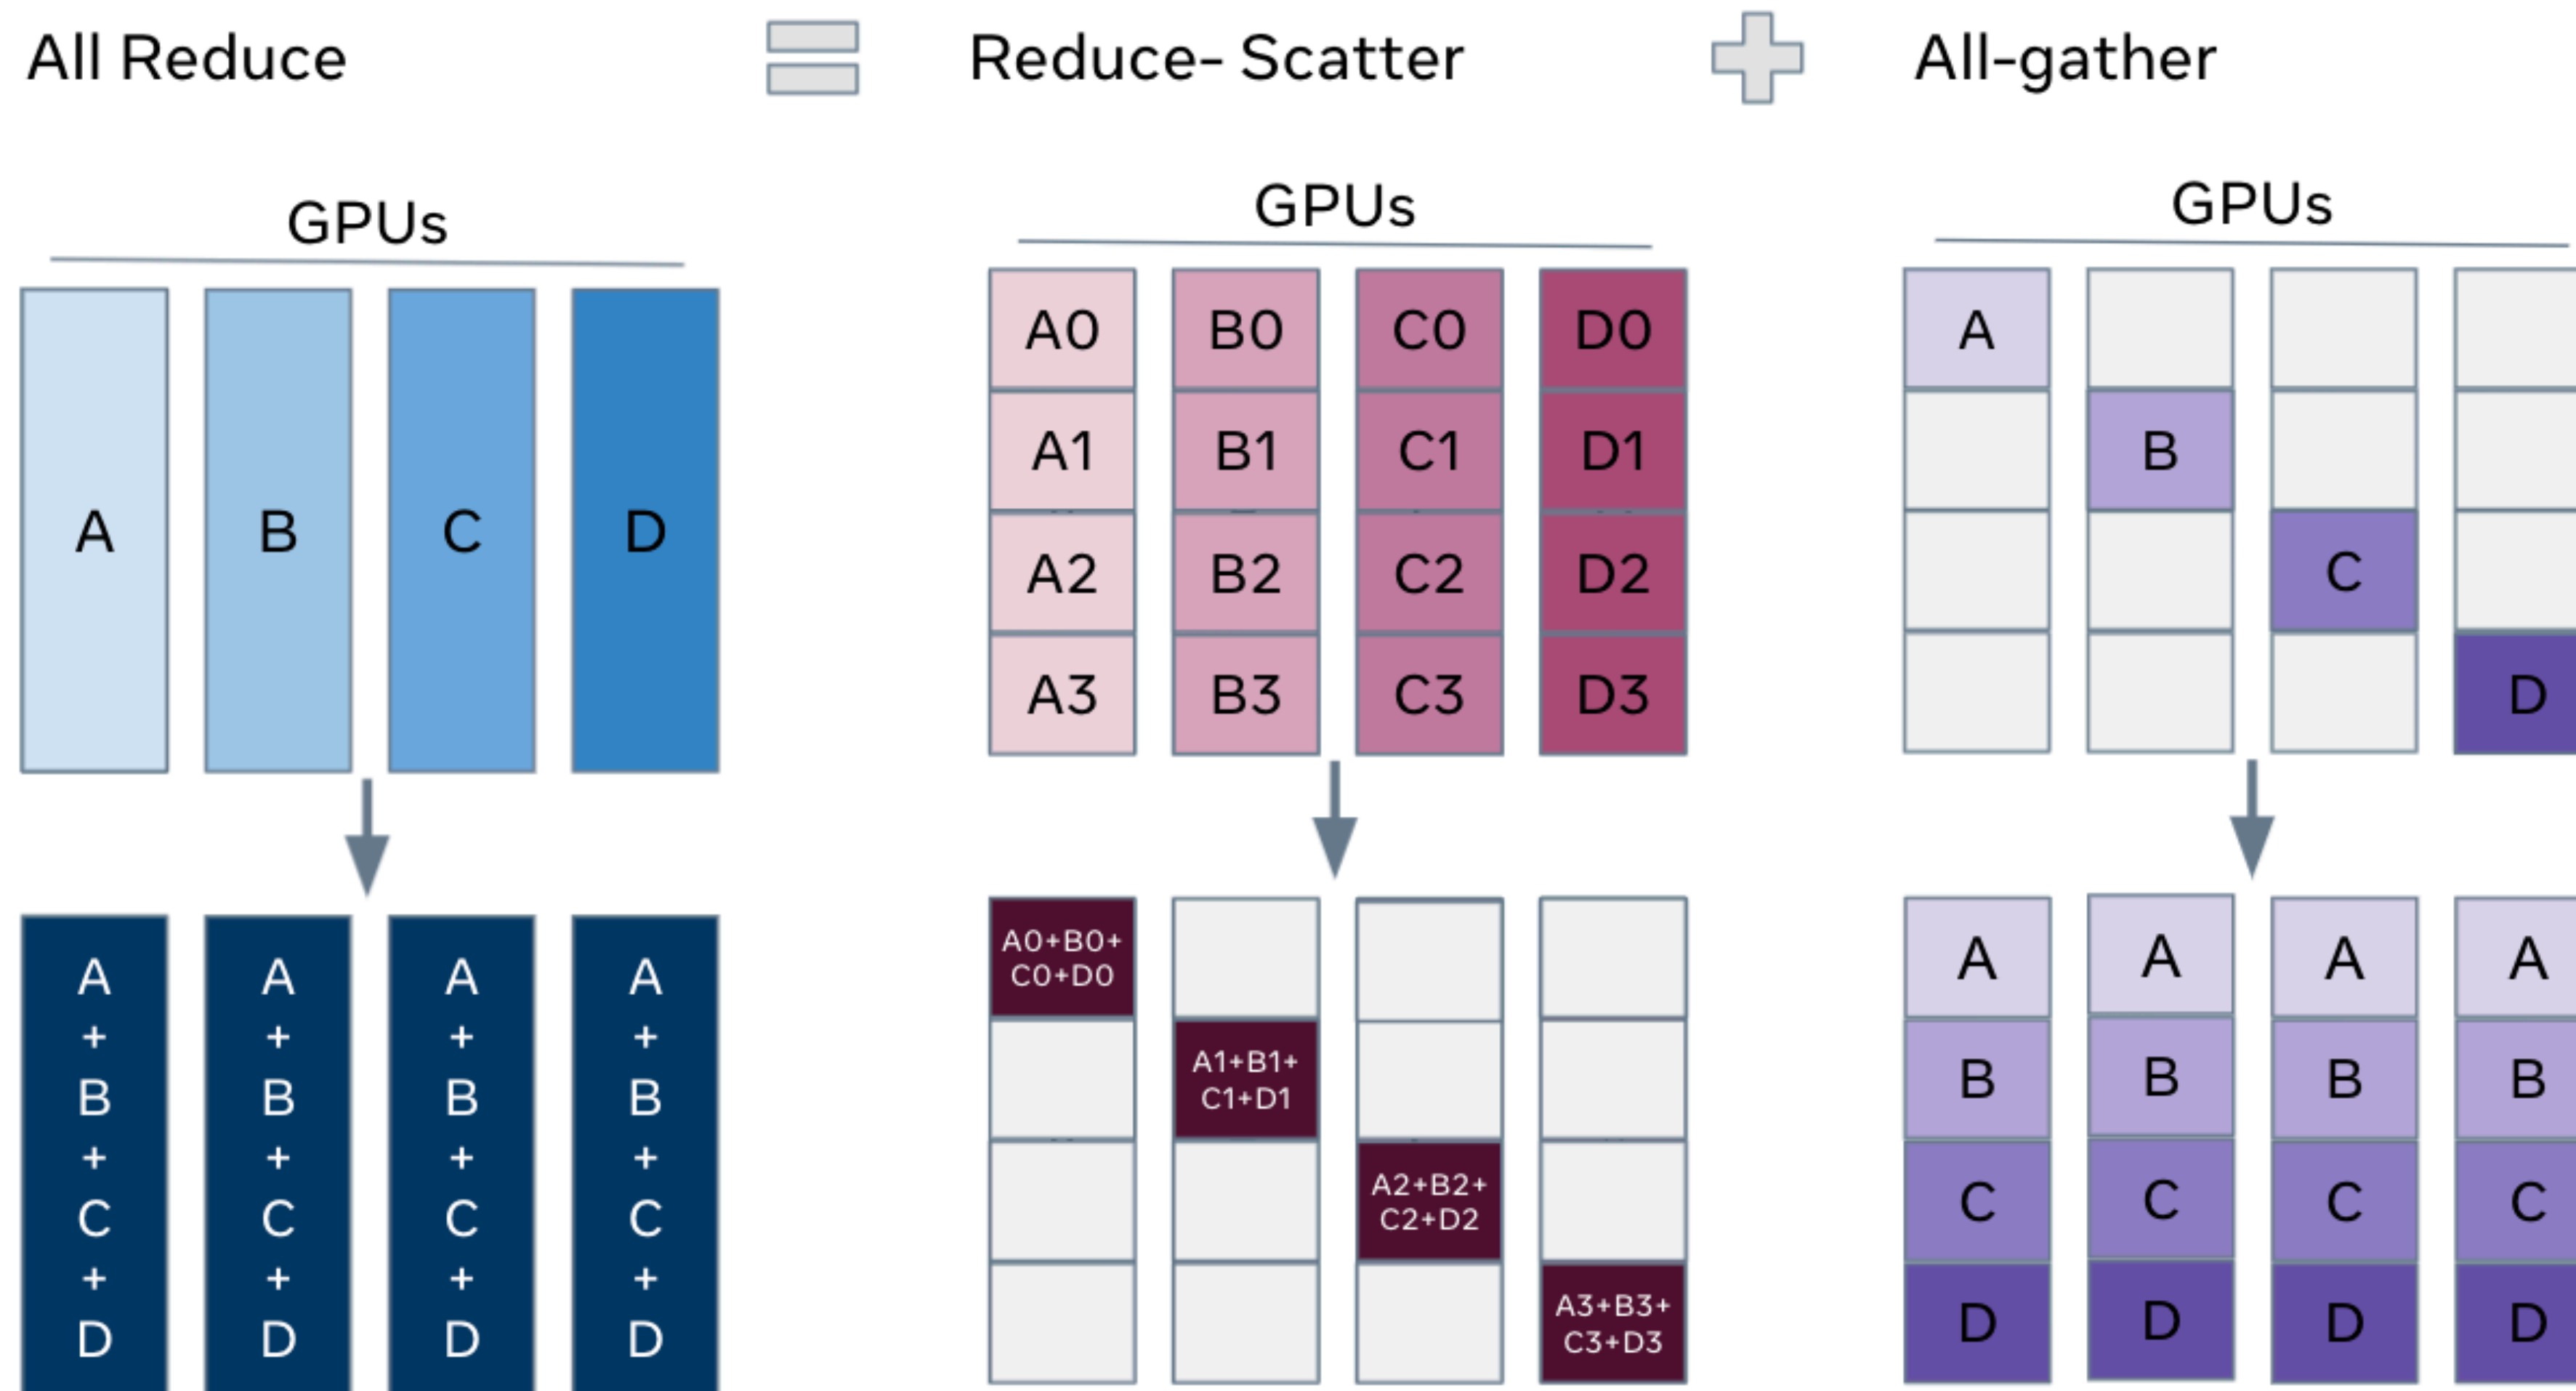


All-reduce



Communication between GPUs


- Some basic communication operations



Distributed Large-Scale Pretraining

- Three Criteria

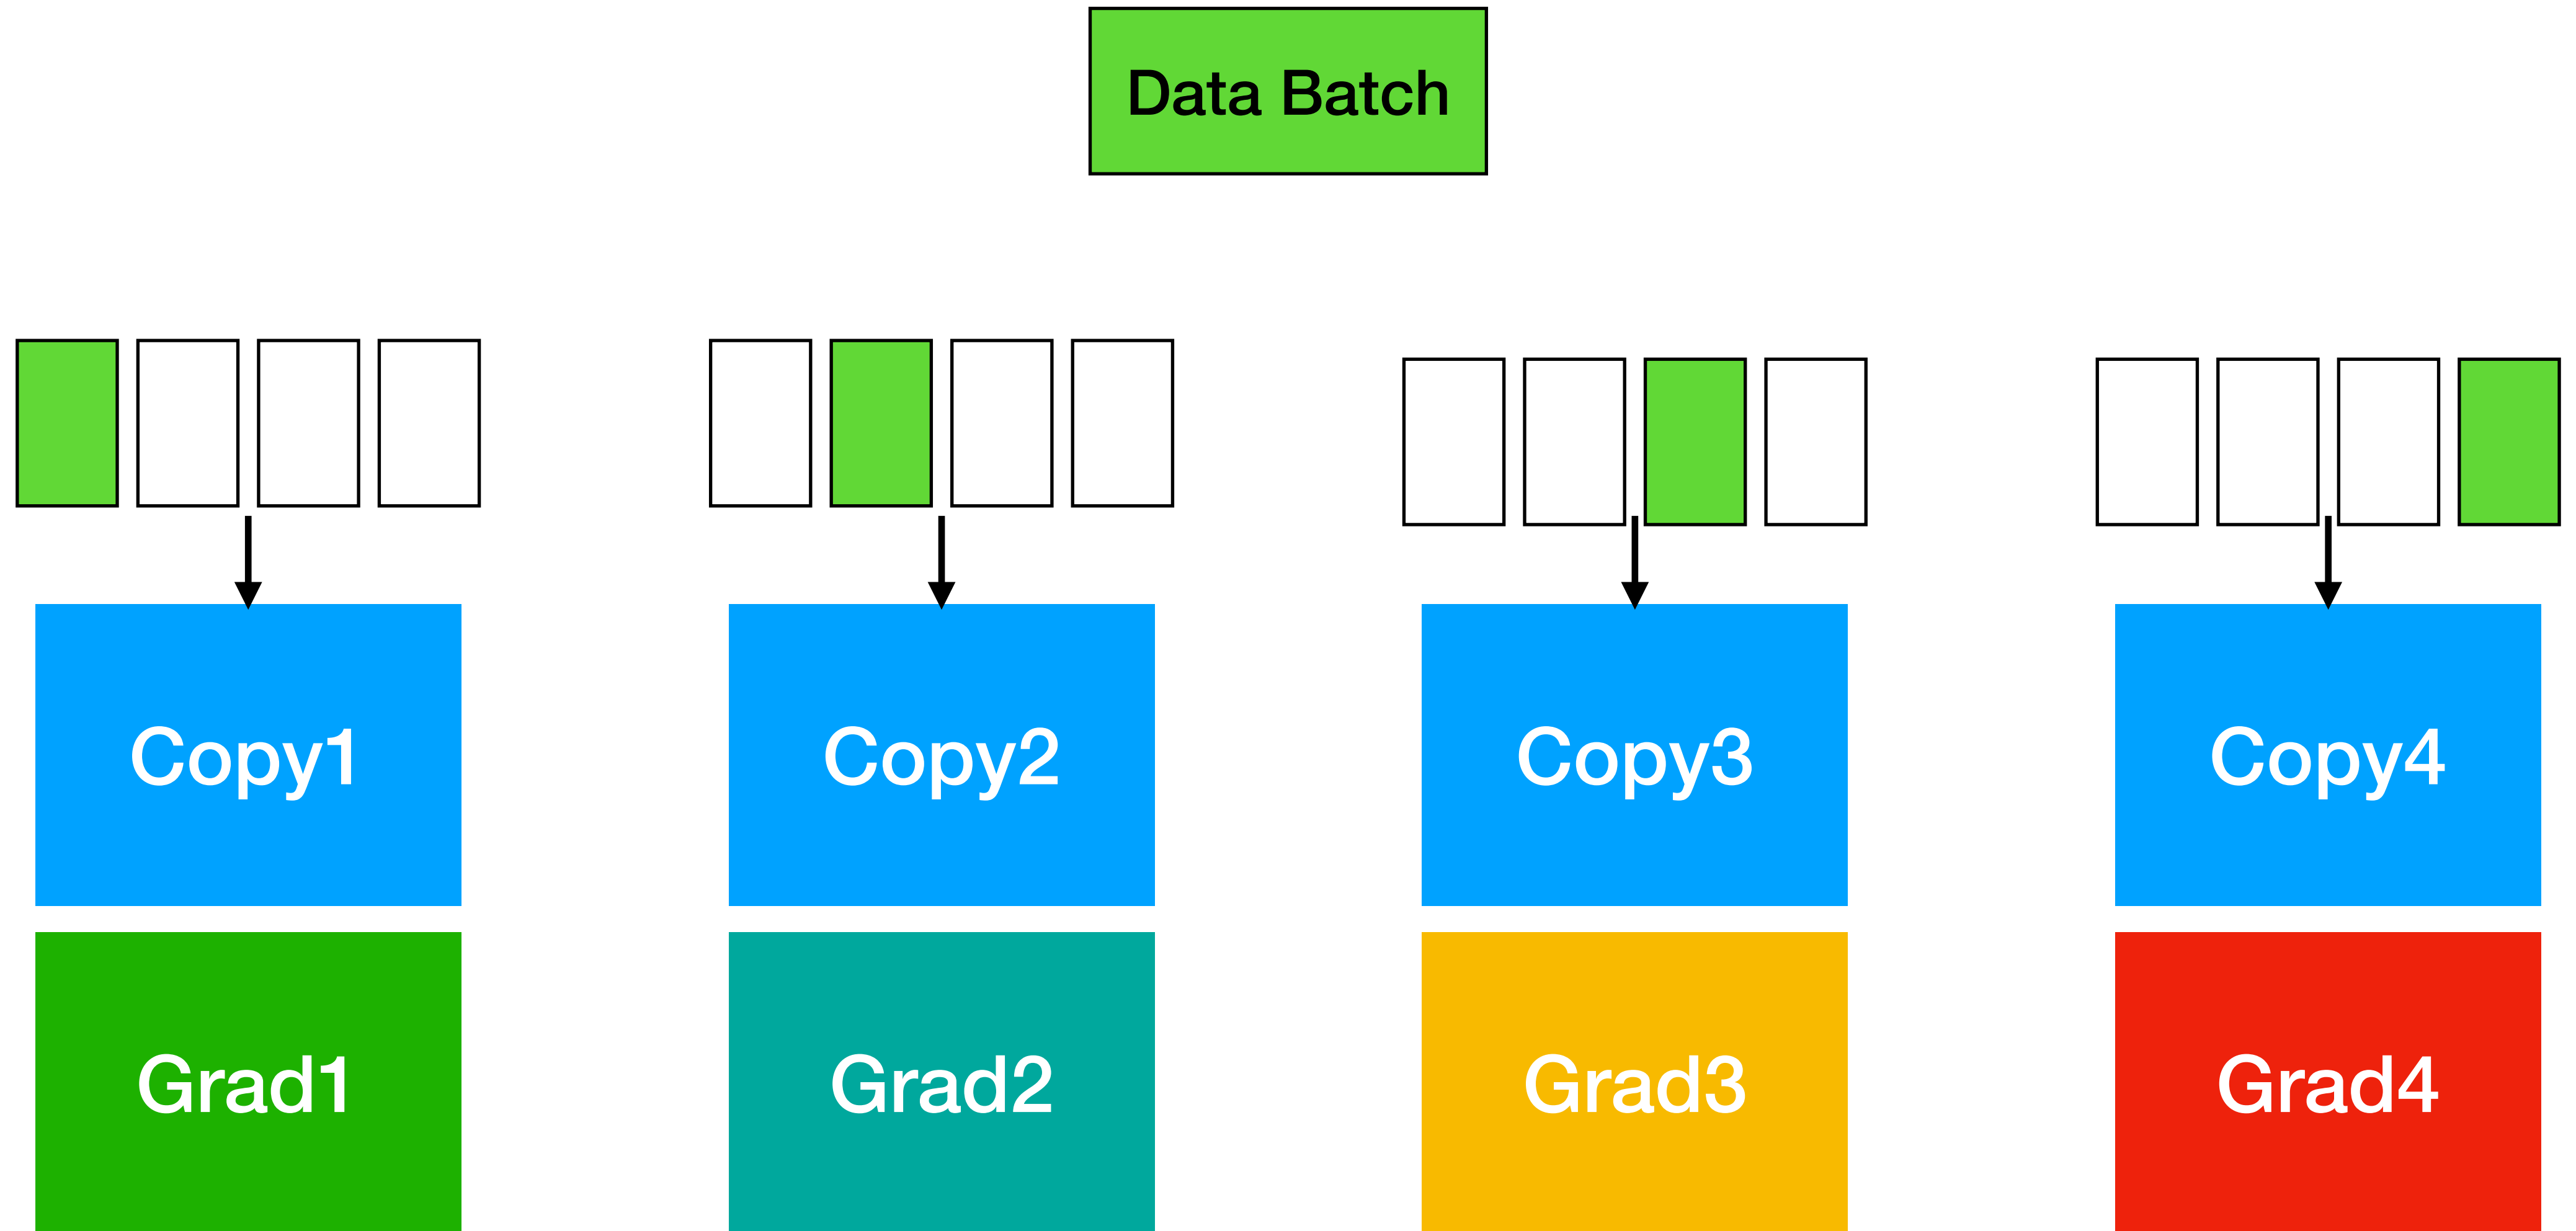
- Minimal **redundant computation**
- Minimal **peak memory cost**
- Minimal **communication overhead**




Distributed Data Parallel (DDP)

- Each GPU keeps a complete copy of the model
- Each GPU process a subset of a training batch
 - No data overlaps between GPUs
- Synchronize parameter gradients after each forward-backward pass

Distributed Data Parallel (DDP)



Distributed Data Parallel (DDP)



$$L = \frac{1}{4}(L_1 + L_2 + L_3 + L_4)$$

Distributed Data Parallel (DDP)

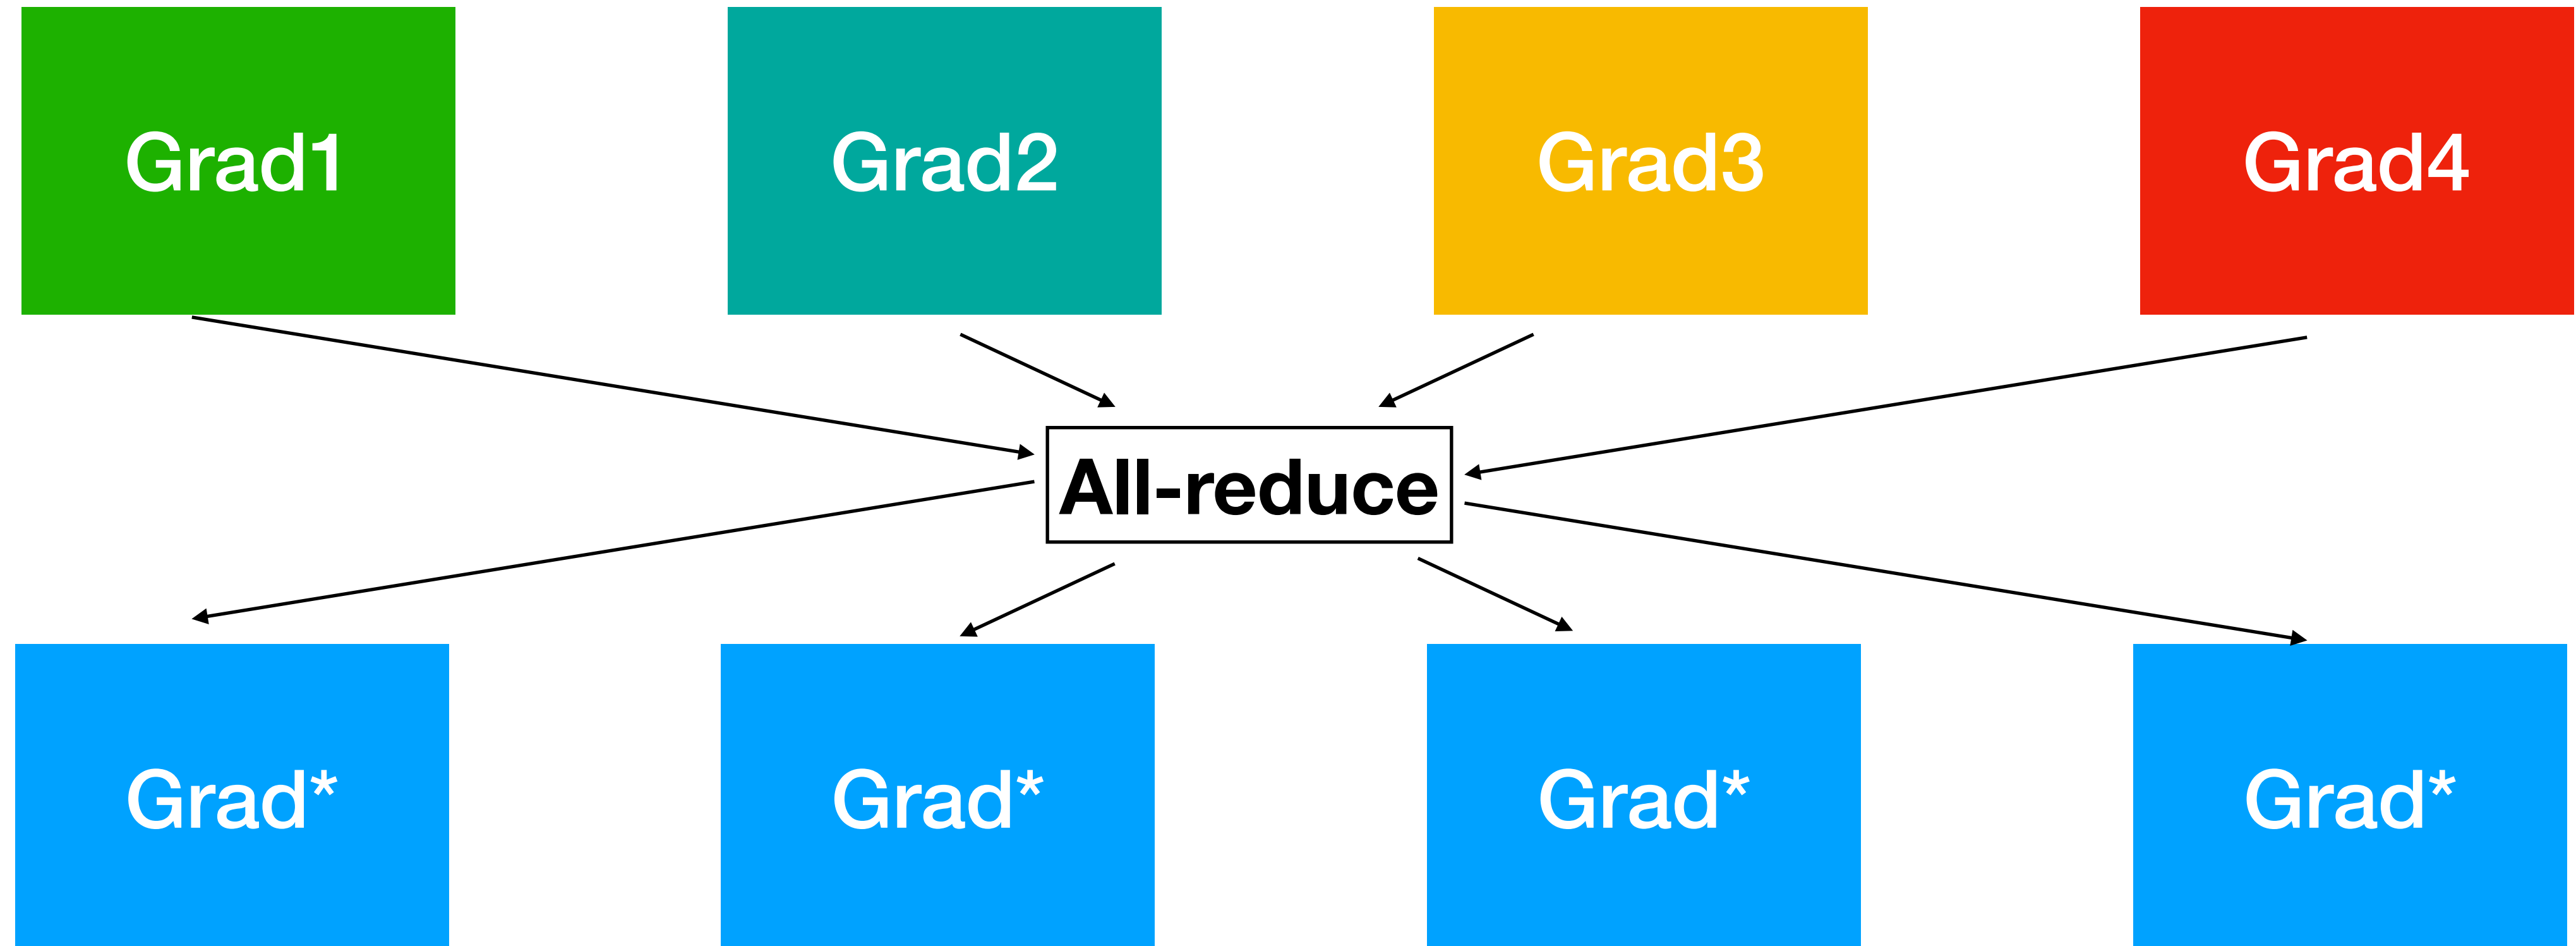
- No redundant computation in each forward-backward pass
- Communication only on gradients
 - One all-reduce operation

Is DDP optimal?

No!

Parameter optimization
is entirely redundant

Distributed Data Parallel (DDP)



$$L = \frac{1}{4}(L_1 + L_2 + L_3 + L_4)$$


Distributed Data Parallel (DDP)

- For large-scale pretraining, parameter optimization **is expensive**
 - A whole copy of 7B model takes **28G** memory
 - Gradient takes the same memory as parameters
 - Optimization states in Adam optimizer take two times of model parameters
 - **28G x 4 = 112G** memory for only storing model and optimization states


How about splitting model parameters together with data?

Distributed Data Parallel (DDP)






Data Batch







All-gather



Distributed Data Parallel (DDP)



Reduce-scatter




Distributed Data Parallel (DDP)

- No redundant computation in optimizing parameters
- No redundant memory cost for the two optimization states in Adam
- No more communication overhead
 - All-reduce = Reduce-Scatter + All-gather
- At one moment, still need to store the whole parameters and gradients

Fully Sharded Data Parallel (FSDP)


- One Transformer model consists of multiple blocks
- Split each block parameters individually
- Gather the parameters of one layer only when we need to compute the output of that layer



Distributed Large-Scale Pretraining


- Three Criteria

- Minimal **redundant computation**
- Minimal **peak memory cost**
- Minimal **communication overhead**




Model/Tensor Parallel


- Specifically designed for linear layers
- Partition both model parameters and hidden states along model dimension
 - Can be applied together with data parallel




FFN




Model/Tensor Parallel



Model/Tensor Parallel




Model/Tensor Parallel



Column Parallel


- Split the weight matrix along the column axis



Column Parallel




Model/Tensor Parallel




Row Parallel

- Split the weight matrix along the row axis



Raw Parallel



$$Y = Y_1 + Y_2$$


Model/Tensor Parallel

- **Communication is heavy**
 - Only applied to GPUs in the **same computing node**
- **How to apply tensor parallel to the attention layer?**
 - First apply column parallel to QKV matrices
 - Compute attentions of different heads in different GPUs
 - Apply row parallel to the output matrix to get the final attention output

Distributed Large-Scale Pretraining

- Three Criteria

- Minimal **redundant computation**
- Minimal **peak memory cost**
- Minimal **communication overhead**



Other Techniques in Large-Scale Pretraining


- Activation Checkpointing

- <https://pytorch.org/docs/stable/checkpoint.html>

- Half-Precision Training


FP32

- 1 bit sign
- 8 bits exponent
- 23 bits fraction
- 1e38 range




FP16

- 1 bit sign
- 5 bits exponent
- 10 bits fraction
- 65504 range



BF16

- 1 bit sign
- 8 bits exponent
- 7 bits fraction
- 1e38 range



Thanks!
Q&A